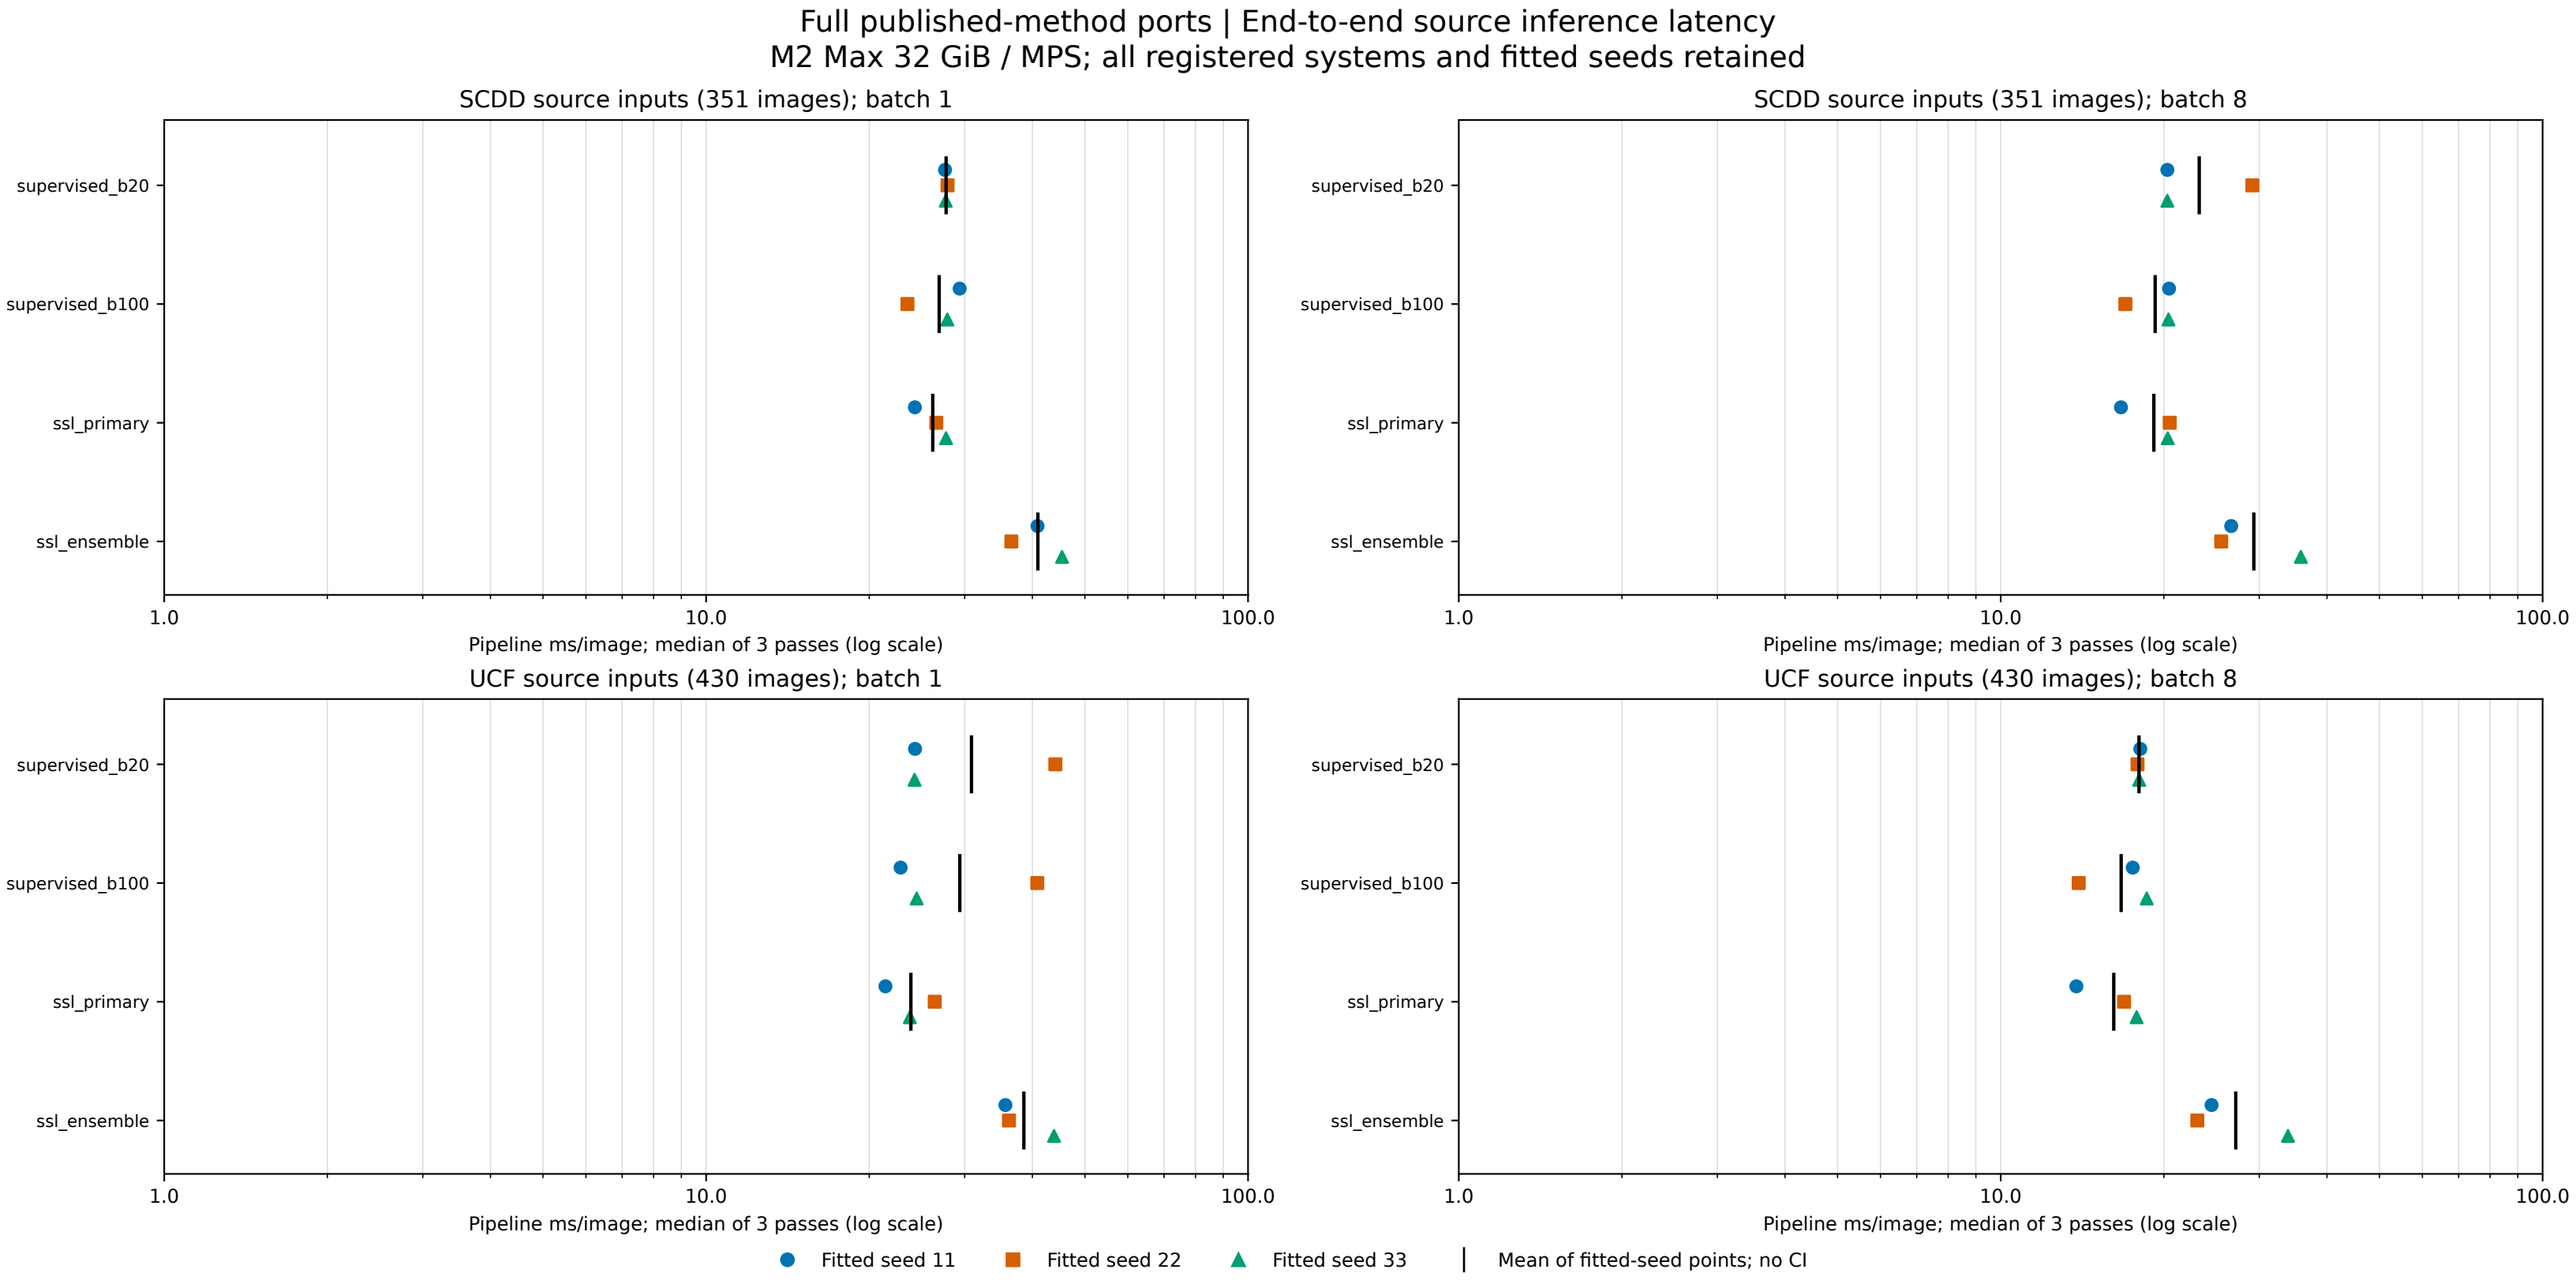
Full published-method ports | End-to-end source inference latency
M2 Max 32 GiB / MPS; all registered systems and fitted seeds retained
SCDD source inputs (351 images); batch 1
SCDD source inputs (351 images); batch 8
UCF source inputs (430 images); batch 1
UCF source inputs (430 images); batch 8
supervised_b20
supervised_b100
ssl_primary
ssl_ensemble
1.0
10.0
100.0
Pipeline ms/image; median of 3 passes (log scale)
Fitted seed 11
Fitted seed 22
Fitted seed 33
Mean of fitted-seed points; no CI